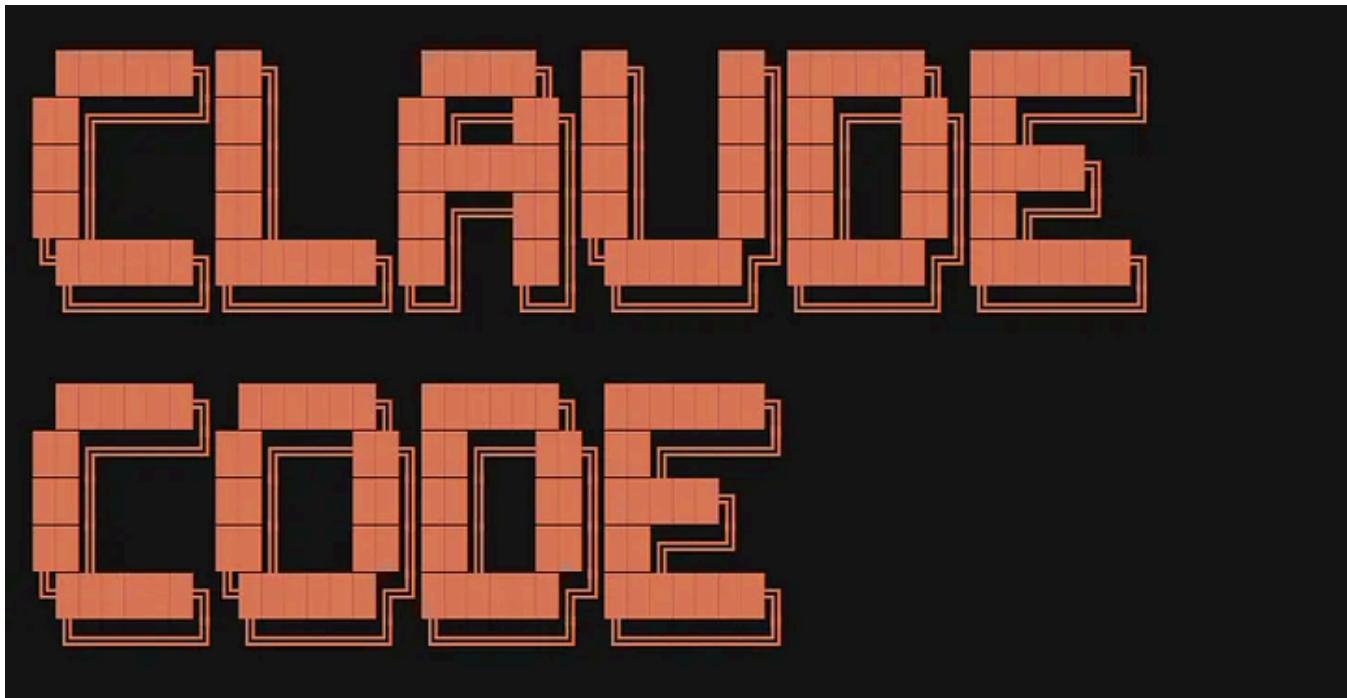


How to Give Claude Code a Memory

Turn your forgetful AI into a reliable coding partner with just four files



Intro

Ever opened Claude Code, only to feel like you’re talking to a brand-new intern who has *no idea* what you’re working on? You patiently explain the project again. Then it creates a file that you already had. Or worse, skips a critical step because it “forgot.”

Claude Code is brilliant but has the memory of a goldfish. 🐠 Each session starts fresh, so context vanishes. The result? Wasted time, duplicated work, and a lot of “Wait, didn’t we already do this?”

But here’s the good news: you can give Claude Code a brain. A simple **four-file framework** stores your project’s memory, so every session feels like you’re working with the *same* developer who knows exactly where things left off.

Let’s see how it works.

Why Claude Code Makes Mistakes

Claude Code isn’t really “aware” of your ongoing project. It doesn’t know what’s been built, what’s still pending, or which tools you’ve chosen.

So when you ask it to keep coding, it might:

- Recreate files it already been made

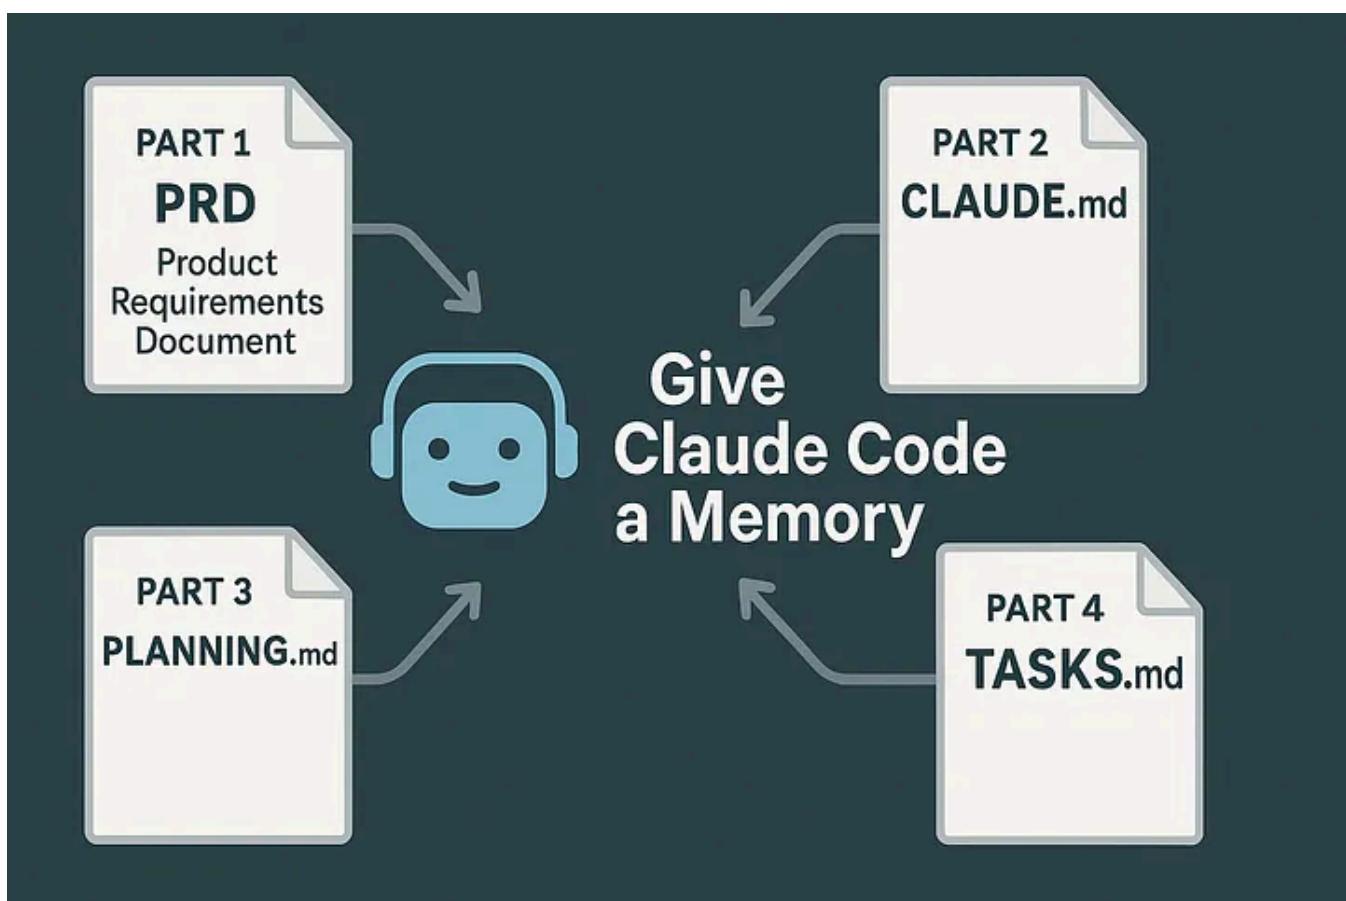
- Miss dependencies or steps
- Forget decisions about tech stack or architecture
- Leave tasks incomplete without realising it

It's like having a helpful assistant who... just keeps forgetting what you said yesterday. The solution? **Persistent external memory** – a place where Claude can “relearn” everything before starting work.

The Simple 4-Part Framework

Here's the fix: you create **four markdown files** that live in your repo. Claude reads these files every session before it touches a single line of code.

- **PRD** → The big picture
- **CLAUDE.md** → Rules for how Claude should behave
- **PLANNING.md** → Your project's architecture and stack
- **TASKS.md** → A living checklist of what to build



Once you have these, Claude *always* knows what you’re building, how you’re building it, and what’s next. No re-explaining. No confusion.

Part 1: PRD (Your North Star)

Begin with a Product Requirements Document — your narrative compass. It answers three anchoring questions: *What are we building? Who is it for? Why does it matter now?* Keep the language plain and executable rather than dreamy. Open with a one-line vision that could sit at the top of a README. Follow with a short problem paragraph that frames the user’s pain in concrete terms, not buzzwords. Describe the core users (two or three succinct personas are plenty), then translate the vision into a handful of well-structured user stories. Add clear success criteria: activation%, response latency targets, retention, maybe a qualitative adoption threshold (“First 20 beta users use chat >3 days in week one”).

Don’t overstuff. The PRD should *inform* the other files, not be copied verbatim into them. Treat it like a stable strategic artefact; when scope shifts, append a dated changelog note at the bottom instead of silently mutating prior intent. That visible evolution helps both you and Claude interpret later architectural shifts as deliberate choices rather than random drift.

Part 2: CLAUDE.md (Behaviour Contract)

Where the PRD states *intent*, CLAUDE.md governs *conduct*. This file tells Claude *how to behave inside this repository*. Think of it as codified ritual plus guardrails. A good CLAUDE.md is concise enough that every directive is salient, yet rich enough that Claude doesn’t improvise unsafe patterns.

Open with a short purpose line (“This file defines session workflow, coding standards, and logging conventions for the Clarity Coach project.”). Then define the invariant session loop in imperative form: load planning context, scan tasks, pick exactly one unblocked atomic task (unless there’s an explicit chain), propose changes, await confirmation, implement, mark completion with timestamp, optionally surface follow-up suggestions. Include a modest section on style: naming conventions, module boundaries, testing expectations (“Each new service function accompanied by a minimal happy-path test”), and a principle like “Favour incremental edits; avoid whole-file rewrites unless structural refactor.” Provide a structured template for session summaries so the log stays uniform over weeks.

Prompt to create CLAUDE.md

Generate a CLAUDE.md file from this PRD that will guide Claude Code sessions on this project.

Mandatory instruction to include inside CLAUDE.md

Always read PLANNING.md at the start of every new conversation, check TASKS.md before starting your work, mark completed tasks to TASKS.md immediately, and add newly discovered tasks to TASKS.md when found.

That single block is the mechanical spine. Without it, Claude reverts to amnesic improvisation. After generation, prune excess adjectives or duplicated rationale — economy here improves compliance.

Parts 3 & 4: PLANNING.md + TASKS.md (Project Memory Core)

PLANNING.md is the distilled architectural worldview. It rephrases (not repeats) the PRD in an implementation lens: current architecture diagram (described textually if no image), stack choices and why (“Postgres over MongoDB for relational integrity and analytic queries”), primary data models, critical request/response flows, latency or security constraints, and any active open questions parked at the bottom. This file should feel *current*; if a decision changes, update *this* first so the next session inherits reality. Resist the temptation to dump exploratory brainstorms here — keep it authoritative, not speculative.

Prompt to create PLANNING.md

Create a PLANNING.md file that includes vision, architecture, technology stack, and required tools list for this app.

TASKS.md is operational oxygen. It encodes momentum as a linear, visible progression. Organize tasks under milestones that reflect value slices (Setup, Data Layer, Retrieval, Chat UI, Testing & Hardening, Deployment). Each task must be atomic enough to be completed in a single Claude iteration: if you’d naturally say “Part 1” or “Phase A,” it’s probably two tasks. Avoid vague verbs (“handle auth”) — instead, say “Implement JWT issuance endpoint (login)”. When a task completes, it is converted in place to `[x]` with a date, optionally adding a brief artefact note (“created `/src/lib/embed.ts`”).

Prompt to create TASKS.md

Create a `TASKS.md` file with bullet point tasks divided into milestones for building this app.

Example fragment (simple is fine):

```
## Milestone 1
- [ ] Initialize repo
- [ ] Set up React + Tailwind
- [ ] Health check endpoint
```

Because Claude will both read and edit this document, clarity of formatting matters more than aesthetic cleverness. Consistency lowers friction.

How to Use It (Daily Loop)

Every working session begins the same way: you instruct Claude to reload the canonical memory (planning + behaviour + operational state) before touching code. This enforces determinism and reduces divergence.

Prompt to initiate building

*Please read `PLANNING.md`, `CLAUDE.md`, and `TASKS.md` to understand the project.
Then complete the first task on `TASKS.md`.*

Claude parses context, selects the top uncompleted task (or you can direct a specific one), proposes changes, implements upon approval, and marks completion. If, during execution, it uncovers missing prerequisites, it should propose and append them rather than silently compensating. That keeps the backlog an honest ledger rather than an after-the-fact fiction.

Before you clear the conversation, snapshot progress into the behavioural log so the next session cold-start fast:

Prompt to add context before clearing

Please add a session summary to `CLAUDE.md` summarizing what we've done so far.

A good summary lists: tasks completed, files created/modified, any key decisions (e.g., “Chose 1024-dim embeddings to reduce token cost”), and proposed next steps. Over time, `CLAUDE.md` becomes the chronological narrative; `PLANNING.md` remains the

architectural present tense; TASKS.md reflects both trajectory (future) and history (checked items); the PRD stays a mostly stable intent artefact.

If something feels off, Claude reimplements a file unnecessarily or skips tests — tighten the directive language in CLAUDE.md or adjust task granularity. Think of the four files as tunable levers: clarity in one compensates for ambiguity elsewhere.

In Short

PRD defines *why and what*. CLAUDE.md enforces *how we work*. PLANNING.md captures *how it's built today*. TASKS.md encodes *what moves next*. Together, they simulate persistent memory and institutional knowledge inside a stateless AI environment. The payoff is fewer repetitions, fewer hallucinated rewrites, cleaner diffs, and a calmer development cadence.

Want a ready-to-commit template bundle or a fully worked sample for a different project type? Just say the word.

Final Thoughts

With just four simple files — PRD, CLAUDE.md, PLANNING.md, and TASKS.md — you turn Claude Code from a forgetful assistant into a reliable partner.

- It always remembers what the project is about
- It never redoess completed work
- It follows the same structure every time
- You have a living roadmap of progress

No more explaining the same thing. No more duplicate files. No more “uhh, what were we building again?” moments.

Next time you start a Claude Code project, set up this mini-memory system first. It takes five minutes and saves *hours* of headaches later.